



# TMS9900 Central Processing Unit

## [Pinout](#)

### Execution speed

#### [Calculations](#)

#### [Examples](#)

#### [Optimizing for speed](#)

#### [Speed test program](#)

### Instruction issues

#### [External instructions](#)

#### [XOP](#)

#### [X](#)

#### [C](#)

#### [Timing diagrams](#)

#### [Electrical characteristics](#)

## Introduction

The TMS9900 is the CPU, i.e. the brain of the TI-99/4A. This microprocessor executes a machine language program located in memory and controls all the other chips in the computer. It's a real 16 bits microprocessor, which means it has 16 data lines and an address space of  $2^{16}$  bytes, i.e. 64K. To accomodate all these lines, TI had to create an extra-large 64-pins chip, that was quite a novelty by that time. They could even afford the luxury of having 5 non-connected pins!

## Pinout

	+-----+-----+-----+	
Vbb	1 o	64 HOLD*
Vcc	2	63 MEMEN*
WAIT	3 T	62 READY
LOAD*	4 M	61 WE*
HOLDA	5 S	60 CRUCLK
RESET*	6	59 Vcc
IAQ	7 9	58 nc
PHI1	8 9	57 nc
PHI2	9 0	56 D15
A14	10 0	55 D14
A13	11	54 D13
A12	12	53 D12
A11	13	52 D11
A10	14	51 D10
A9	15	50 D9
A8	16	49 D8
A7	17	48 D7
A6	18	47 D6
A5	19	46 D5
A4	20	45 D4
A3	21	44 D3
A2	22	43 D2
A1	23	42 D1
A0	24	41 D0
PHI4	25	40 Vss
Vss	26	39 nc

Vdd	27	38	nc
PHI3	28	37	nc
DBIN	29	36	IC0
CRUOUT	30	35	IC1
CRUIN	31	34	IC2
INTREQ*	32	33	IC3
+-----+			

### Power supply

**Vbb:** -5V Vcc: +5V Both pins (2 and 59) must be connected.

**Vdd:** +12V Vss: Ground. Both pins (26 and 40) must be connected.

### Clock

**Phi1-Phi3** are 4 input pins that receive the same signal from the TMS9904 clock generator, with one exception: each signal is shifted by 1/4 of a phase with respect the the previous one.

### Data bus

**D0-D15:** These 16 pins are used to read or write data. Note that contrarily to almost anybody else, TI made D0 the most significant bit (weight >8000) and D15 the least significant bit (weight >0001).

### Address bus

**A0-A14:** These 15 output pins are used the specify the address of 32Kwords in memory. Each word is two bytes long (since we have 16 data lines), thus we are effectively addressing 64Kbytes. On the TI-99/4A, the data bus is multiplexed as 2 x 8 bits for almost every purpose except accessing the console ROMs and the scratch-pad RAM. The multiplexing circuitry controls a pseudo-address line A15 that indicated whether the 8-bit data bus contains the most significant byte of the 16-bit bus (A15=0) or the least significant byte (A15=1).

### Bus control

**MEMEN\*** Memory enable. When active low, this pin indicated that the TMS9900 wants to access memory and has placed a valid address on the address bus.

**DBIN** Data bus in. When active (high) this pin indicates that the TMS9900 is ready to accept data.

**WE\*** Write enable. When active (low) this pin indicates that the TMS9900 has placed valid data on the bus.

### CRU control

**CRUCLK** CRU clock. When active (high) this pin indicates the the TMS9900 is performing a CRU operation (or an external instruction).

**CRUOUT** This output pin contains the data that the TMS9900 sends out during CRU operations.

**CRUIN** This pin is used by the TMS9900 to input data during CRU operations.

### Interrupt control

**INTREQ\*** When active (low) this input pin signals the TMS9900 that an interrupt is pending. If it accepts the interrupt, the TMS9900 will perform BLWP @>0000 through BLWP @>003C depending on the interrupt level.

**IC0-IC3** These 4 input pins indicate the level of the interrupt (0-15). The LIM1 instruction can be used to define the "cutoff" level, above which the TMS9900 will ignore interrupts. On the TI-99/4A those pins are hardwired as low,low,low,high which means all interrupts are level 1.

**LOAD\*** Non-maskable interrupt. When active low, this pin forces the TMS9900 to perform a BLWP @>FFFC interrupt. If it remains low, interrupts continue to be issued, thus is should not remain low for more than one instruction.

**IAQ** Instruction acquisition. This output pin indicates that the TMS9900 is acquiring an instruction. It can be used to detect illegal opcodes or to prevent LOAD\* to last longer than one instruction. On the TI-99/4A, IAQ

and HOLDA are combined via an OR gate and presented to the peripheral port. However, the flex cable connector does not carry that signal to the PE-Box.

**RESET\*** When active (low) this input pin resets the TMS9900 and inhibits WE\* and CRUCLK. RESET\* must remain low for at least 3 clock cycles. As soon as it becomes high again, the TMS9900 performs a BLWP @>0000 (note that it is the same vector as interrupt level 0).

#### Memory control

**HOLD\*** When active (low) this pins tells the TMS9900 that a DMA controller wants to perform Direct Memory Access. The TMS9900 set D0-D15, A0-A14, MEMEN\*, DBIN and WE\* in high impedance state (isolated) then activate HOLDA and waits until HOLD\* becomes high again. On the TI-99/4A this pin is hardwired high, which means we cannot perform DMA operations.

**HOLDA** Hold acknowledge. This output pin is used to tell the DMA controller that the it can perform direct memory access.

**READY** When active (high) this input pin tells the TMS9900 that the memory is ready to read or write data. If it's low, the TMS9900 enters a wait state and suspends all operations until READY becomes high again. On the TI-99/4A this line is used to multiplex the data bus, i.e. handle it as two 8-bit bytes, instead of one 16-bit word. **WAIT** When active (high) this output pin indicates that the TMS9900 is now in wait state.

## Execution speed

To execute an instruction, the TMS9900 must first fetch it from memory, which takes a few clock cycles (depending on the memory), then it must execute the instruction which takes more clock cycles (depending on the instruction) and may require fetching one or two arguments from the memory (again, more clock cycles according to the addressing mode and memory type).

On the TI-99/4A, there are two kinds of memory: 16-bit and 8-bit. Console ROMs (address >0000-1FFF) and the RAM scratch-pad (address >8300-83FF) are the only 16-bit memories. All the rest, including peripheral cards, and memory-mapped devices (GROM, VDP, sound and speech chips) are accessed in a byte-wise manner. This requires multiplexing the 16-bit data bus in two 8-bits chunks. An electronic circuitry in the console takes care of that burden and uses the READY line to halt the TMS9900 until the peripheral has received/sent the second data byte. Therefore, accessing such a memory results in 4 wait states for each memory access.

The table below can be used to calculate how long an instruction takes to be executed.

- The first column lists the instructions in alphabetical order.
- The second column indicates how many clock cycles are required to execute that instruction in its most primitive form (generally: using workspace registers).
- The third column indicates how many memory access operations are required to fetch the instruction from the program memory and the operands from their memory. In most cases, only one memory access is needed to retrieve the instruction, except for immediate instructions that need a second cycle to retrieve the immediate value embedded in the program. If the program is in slow memory, it will take 4 more clock cycles for each of these. All other memory accesses are used to read and write the arguments. If the memory in question is not the console ROMs nor the scratch-pad RAM, add 4 clock cycles per access operation.
- The last two columns indicate whether the instruction required fetching arguments. Some don't, some only need a source arguments, some need two arguments (source and destination).
- In case the arguments are not register, the instruction may require more clock cycles and/or memory access than what's listed in column 2 and 3. Table 2 allows to calculate the number of extra clock cycles required to access arguments.

Instruction	Clock cycles	Memory access	Source	Destination
A	14	4	Y	Y

AB	14	4	Y	Y
ABS(pos) (neg)	12	2	Y	-
	14	3	Y	-
AI	14	4	-	-
ANDI	14	4	-	-
B	8	2	Y	-
BL	12	3	Y	-
BLWP	26	6	Y	-
C	14	3	Y	Y
CB	14	3	Y	Y
CI	14	3	-	-
CKOF	12	1	-	-
CKON	12	1	-	-
CLR	10	3	Y	-
COC	14	3	Y	-
CZC	14	3	Y	-
DEC	10	3	Y	-
DECT	10	3	Y	-
DIV (ovf) (no ovf)	16	3	Y	-
	92-124 (1)	6	Y	-
IDLE	12	1	-	-
INC	10	3	Y	-
INCT	10	3	Y	-
INV	10	3	Y	-
Jump (taken) (not taken)	10	1	-	-
	8	1	-	-
LDCCR	20 +2*bits	3	Y	-
LI	12	3	-	-
LIMI	16	2	-	-
LREX	12	1	-	-
LWPI	10	2	-	-
MOV	14	4	Y	Y
MOVB	14	4	Y	Y
MPY	52	5	Y	-
NEG	12	3	Y	-
ORI	14	4	-	-
RSET	12	1	-	-
RTWP	14	4	-	-
S	14	4	Y	Y
SB	14	4	Y	Y
SBO	12	2	-	-
SBZ	12	2	-	-
SETO	10	3	Y	-

Shift	12 +2*disp	3	-	-
(disp in R0)	20 +2*disp	4	-	-
SOC	14	4	Y	Y
SOCB	14	4	Y	Y
STCR (1-7)	42	4	Y	-
(8 bits)	44	4	Y	-
(9-15 bits)	58	4	Y	-
(16 bits)	60	4	Y	-
STST	8	2	-	-
STWP	8	2	-	-
SWPB	10	3	Y	-
SZC	14	4	Y	Y
SZCB	14	4	Y	Y
TB	12	2	-	-
X (note 2)	8	2	Y	-
XOP	36	8	Y	-
XOR	14	4	Y	-
Illegal	6	1	-	-
Interrupts	22	5	-	-
Reset	26	5	-	-

### Notes

- 1) DIV execution time, when no overflow occurs, depends on the partial quotient after each clock cycle during execution.
- 2) For X, add this time to the execution time of the instruction found at the source address, minus 4 clock cycles and 1 memory access.

For each source and destination arguments (if any) add the following:

Address mode	Clock cycles	Memory access
Rx	0	0
*Rx	4	1
*Rx+ (byte)	6	2
(word)	8	2
@>xxxx	8	1
@>xxxx(Rx)	8	2

### Note

For the \*Rx+ addressing mode, the number of clock cycles depends on whether the register must be incremented by 1 or by 2. The byte-oriented operations increment it by 1 and use 6 clock cycles, these are: AB, CB, MOVb, SB, SOCB and SZCB. In addition, the LDCR and STCR are considered as byte operations if they transfer 1 to 8 bits (with 9 to 16 bits they are word operations and use 8 clock cycles).

## Examples of calculations

### LIMI 2

The LIM2 instruction uses 16 clock cycles and 2 memory access operations: one to fetch the LIM2 instruction, one to fetch the immediate value 2.

Execution: 16 cycles

```
1 mem access to read LIM2 }
1 mem access to read 2    } add 8 cycles if program is in slow mem.
```

This adds up to:

16 cycles if the instruction is in the ROMs or the scratch-pad,

$16 + 2 \times 4 = 24$  clock cycles otherwise (remember, there are 4 wait states per memory access).

With a 3 MHz clock that cycles every 333 nanoseconds, this boils down to  $16 \times 333 \text{ ns} = 5.33 \text{ microseconds}$  or  $24 \times 333 \text{ ns} = 8 \text{ microseconds}$ , depending on which memory the instruction is in. Tip: divide the number of clock cycles by 3 to find the execution time in microseconds (with a 3 MHz clock).

### CLR R2

The CLR instruction uses 10 clock cycles and 3 memory access operations. Depending on which memory it is in, it may require from 10 to 22 clock cycles. But CLR also takes an argument that we must consider. In this case, the argument is a register and the clock cycles needed for its access can be found in table 1.

Execution: 10 cycles

```
1 mem access to read CLR: add 4 cycles if program is in slow mem.
1 mem access to read R2   }
1 mem access to write R2  } Add 8 cycles if workspace is in slow mem
```

### CLR @TEST

The CLR instruction itself still uses 10 or 22 clock cycles to execute, but now dealing with the argument requires 8 extra clock cycles and 1 extra memory access operations (to read the address of TEST from the word following CLR in the program).

Execution: 10 cycles

Argument @xxx: 8 cycles

```
1 mem access to read CLR   }
1 mem access to read TEST  } Add 8 cycles if program is in slow mem
1 mem access to read the contents of @TEST }
1 mem access to write to @TEST           } Add 8 cycles if TEST is in slow mem
```

If the address TEST is not in the ROMs nor in the scratch-pad, it will add  $8 + 2 \times 4 = 16$  clock cycles to the execution time. Otherwise it just adds 8 cycles. We thus have:

```
10+8    = 18 cycles if program and TEST are both in fast memory.
10+8+8  = 26 cycles if program in slow memory, but TEST is in fast memory.
10+8+8  = 26 cycles if program is in fast memory, but TEST is in slow memory.
10+8+8+8 = 34 cycles if both program and TEST are in slow memory.
```

### CLR \*R2

Fetching the source argument from the address found in register R2 requires 4 clock cycles and 1 memory access operation. Depending whether the workspace is in the scratch-pad or not (no workspace should ever be in ROM!), this may require 4 additional clock cycles. Depending whether the target address found in R2 (i.e. \*R2) is in slow or fast memory, we may have to add 4 more cycles.

Execution: 10 cycles

Argument \*Rx: 4 cycles

```
1 mem access to read CLR: Add 4 cycles if program is in slow mem
1 mem access to read R2: Add 4 cycles if workspace is in slow mem
1 mem access to read the contents of *R2 }
1 mem access to write to *R2           } Add 8 cycles if R2 points to slow mem
```

So the total could be:

$10+4 = 14$  cycles if CLR, workspace and \*R2 are all in fast memory.  
 $10+4+4 = 18$  if either CLR or the workspace is in slow memory.  
 $10+4+4+4 = 22$  if both CLR and workspace are in slow memory  
 $10+4+8 = 22$  if R2 points to slow memory  
 $10+4+8+4 = 26$  if R2 points to slow memory and either CLR or WS is in slow memory  
 $10+4+8+4+4 = 30$  if everything is in slow memory.

### **LDCR R1,7**

The LDCR instruction itself requires 20 clock cycles. There may be additional cycles required to retrieve the instruction, or the register. Writing the CRU definitely implies 2 additional clock cycles per bit transferred: in this case we are transferring 7 bits, thus we'll eat  $2*7=14$  cycles.

Execution: 20 cycles

1 mem access to read LDCR: Add 4 cycles if program is in slow mem  
 1 mem access to read R1: Add 4 cycles if workspace is in slow mem  
 1 cru write operation: Add 2 cycles per bit transferred (here:  $2*7=14$ )

### **LDCR \*R1+,7**

Now this one gets tricky: the LDCR instruction requires  $20+2*7=14$  clock cycles plus whatever may be needed for the memory access operations, just as above. But now we must allocate time to increment R1 after execution. Since we are transferring only one byte (that is, less than one byte: only 7 bits), R1 will be incremented by 1, which requires 6 clock cycles and 2 memory accesses. Depending on where the workspace is, this could add up to  $6+2*4=14$  cycles.

Execution: 20 cycles

1 mem access to read LDCR: Add 4 cycles if program is in slow mem  
 1 mem access to read R1: Add 4 cycles if workspace is in slow mem  
 1 mem access to read \*R1: Add 4 cycles if source is in slow mem  
 Incrementing R1: 6 cycles (byte operation)  
 1 mem access to write back R1: Add 4 cycles if workspace is in slow mem  
 1 cru write operation: Add 2 cycles per bit transferred (here:  $7*2=14$ )

### **LDCR \*R1+,14**

Here, we are transferring 14 bits, thus the LDCR instruction takes  $20+2*14=48$  cycles, plus what's needed for the memory access. In addition, since we are transferring more than 1 byte, R1 will be incremented by two, which requires 8 clock cycles instead of the 6 mentioned above.

Execution: 20 cycles

1 mem access to read LDCR: Add 4 cycles if program is in slow mem  
 1 mem access to read R1: Add 4 cycles if workspace is in slow mem  
 1 mem access to read \*R1: Add 4 cycles if source is in slow mem  
 Incrementing R1 by two: 8 cycles (word operation)  
 1 mem access to write R1: Add 4 cycles if workspace is in slow mem  
 1 cru write operation: Add 2 cycles per bit transferred (here:  $14*2=28$  cycles)

### **SRL R2,4**

The SRL shift operation requires 12 clock cycles, plus 2 cycles for each position shifted. Since the displacement is 4 in this example, it will require  $12+2*4=20$  cycles. Plus the number of cycles required for 3 memory access operations: 0 to 12 depending on the position of the program and the workspace.

Execution: 12 cycles

Shifting: 2 cycles per position

1 mem access to read SRL: Add 4 cycles if program is in slow mem  
 1 mem access to read R2   }  
 1 mem access to write R2   } Add 8 cycles if workspace is in slow mem

### **SRL R2,0**

Here, we are fetching the displacement from R0 (let's say it contains 5). This indirect shift operation requires  $20+2*5=30$  cycles, and 4 memory access operations instead of 3, as we must now fetch R0.

Execution: 12 cycles  
 Shifting: 2 cycles per position  
 1 mem access to read SRL: Add 4 cycles if program is in slow mem  
 1 mem access to read R2 }  
 1 mem access to read R0 } Add 12 cycles if workspace is in slow mem  
 1 mem access to write R2 }

These calculations are a pain to perform, aren't they? I'm playing with the idea to write an optimisation helper, i.e. a program that would read an assembly source file, and produce a corresponding output file with the execution times listed as comments. But I don't know when I will have time for that. [Anyone](#) aware of such a program around here?

## Optimizing for speed

Now we can see what are the cycle-hungry operations: DIV, MPY, LDCR, STCR, XOP and BLWP.

That's why it is often wise to perform a multiplication using shifts and additions rather than MPY:

```
MPY  R0,R8
```

Requires 72 cycles to execute (52 in 16-bits memory). And that's the fastest MPY.

Now if R0 contains 8, we could have written:

```
SLL  R8,3
```

Which does the same, but only uses 30 cycles (18 in fast memory).

To multiply by ten, we could do:

SLL  R8,1	Multiply by two
MOV  R8,R9	
SLL  R8,2	And then by 4 (which makes 8)
A    R9,R8	Add it up: 2+8=10.

This requires 58 cycles in fast memory and 114 in slow memory. True, this is slower than the initial MPY, but we may have a use for the intermediary result in R9 (that is, R8 times two).

For the same reason, many programmers avoid calling subroutines with BLWP-RTWP and favor BL-B \*R11, at least in critical regions of their programs.

We can't do much about LDCR and STCR, but this is less of a problem: these instructions are rarely used anyhow, and the limiting factor may well be the hardware they are addressing (although not very likely: TTLs are fast).

## Test program

All this theory is impressive, but we'd like to verify whether it is true in the real world. Let's write a little test program and time its execution with a stop watch (you may want to time it automatically, with the [TMS9901 timer](#), but that's another story).

START	LWPI >A800	Load our workspace	** 1 **
	LI R1,DELAY	The subroutine we want to time	
	LI R2,>B000	Where it will run	** 2 **
	MOV R2,R3		
	LI R0,EOPG-DELAY	Subroutine size	
LP0	MOV *R1+,*R2+	Copy subroutine in target memory	
	DECT R0		



```

*      JNE  LP0
      B    @RET          To return immediately  ** 3 **
      B    *R3          Call subroutine

DELAY  LI    R1,100      You can change this value
LP1    CLR   R0
LP2    DEC   R0          Inside loop, executes 65536 times
      JNE   LP2
      DEC   R1          Outside loop
      JNE   LP1

RET    LWPI  >20BA      Assuming editor/assembler workspace
      B    *R11         Done, return to editor/assembler module
EOPG   END

```

The test program first copies the timed routine in memory. This could be the scratch pad memory or the memory expansion. Then it executes the delay loop. Once it is done, it returns to the caller. I have assumed that this is the Editor/Assembler cartridge (oe Funnelweb). If it's not, modify the return instructions accordingly.

Now let's do some measurements. First of all, assemble the program with the B @RET in line \*\* 3 \*\*. The delay loops will be skipped and the program will return immediately. This allows us to account for the time it takes the Editor/Assembler module to enter our program, and to display the <press any key> message when returning from it. As you'll see, this is so fast that we cannot time it..

Then let's comment out line \*\* 3 \*\* and time our program in four different situations: Modify line \*\* 1 \*\* to use a workspace in the memory expansion (>A800) or in the scratch-pad (>83E0). For each of those, modify line \*\* 2 \*\* to copy the delay loop in the memory expansion (>B000) or in the scratch-pad (>8300). Write the resulting times in this table:

<b>Program</b>	<b>Memory expansion</b>	<b>Scratch-pad</b>
Workspace		
Memory expansion	89 sec (100%)	62 sec (70%)
Scratch-pad	62 sec (70%)	44 sec (49%)

Of course, the faster way is to have both the program and the workspace in the scratch-pad: in our case, it's twice as fast as the slowest solution. Unfortunately, this is not practical as the scratch-pad is only 256 bytes long. And many of those bytes have special meanings. Most of the time however, it is possible for you to place your workspace in the scratch-pad: LWPI >8300 for instance. This will substantially speed up your program (in our case, by 30%), especially if you are carefull to reserve your registers for frequently used variables.

Now, if there are some speed-critical routines in your program (such a scrolling the screen left/right in an arcade game), and if they are small enough, you could copy them in the scratch-pad as we did above and execute them there. Say at >8320, not to overwrite your own workspace.

But anyhow, there is more to gain with the workspace than with the program. Consider the instruction A R1,R2 for instance: it requires one memory operation to ftech the instruction from the program memory, and three accesses to the workspace (to read R1, to read R2 and to write back R2). So by placing the workspace in fast memory you gain three times more than by plancing the instruction itself in fast memory.

To speed up the program itself, you could optimize your code to avoid using those instructions that require a lot of time to execute, as discussed above. And often an improved algorithm will do better than nay optimization trick!

## Instructions issues

## External instructions

- We could use the CKON\* and CKOF\* lines to switch the clock speed, by feeding the appropriate signals to the TMS9904.
- More generally, we could use those line to activate any kind of hardware we want.

For instance, I was told (by Anders Persson) that the "Cortex" computer, which was sold as a kit in the 80s, made use of these instructions for the following:

- RSET caused a hardware reset
- IDLE lit an external LED. The system would issues lots of IDLE when it had nothing to do, so you could see that.
- CKON and CKOF enabled and disabled a memory mapper system, in the style of the one used by the SuperAMS card.
- LREX triggered a delayed interrupt after two instructions (via a series of flip-flops clocked by the IAQ line). This was used by a debugger to execute single instructions, by setting up a pseudo-RTWP pointing at a given instruction, and then doing LREX RTWP. The RTWP would go back to the desired instruction, which was executed, and then the interrupt would fire and the debugger was back in control. Nifty, isn't it?

## XOP instruction

The XOP (eXtended OPeration) instruction is kind of a special BLWP. It takes a source argument and an operation number from 0 to 15. The XOP instructions uses that number to perform a BLWP to one among 16 vectors located in memory addresses >0040-007F. In addition, it places the content of the source argument in the R11 register of the new workspace. Finally, bit 6 (weight >0200) is set in the status register while a XOP instruction is executed. Note that the TMS9900 does not test the INTREQ\* interrupt request pin after a XOP operation.

The main advantage of XOP is that it only requires one word. Therefore, we could use it to replace any word in a program and interrupt its execution. That's how my debugger RIP v.2 works: to set a breakpoint it saves the content of a memory address and replaces it with an XOP 1. Execution of this XOP results in activating RIP that can then execute the saved instruction and/or ask the user what to do. True, BLWP \*Rx is only one word long, and so is BLWP Rx, but both expect special values in the registers. The first one want a pointer to the WR and PC vectors in Rx, the second want the WR vector in Rx and the PC vector in the next register. XOP is much more convenient.

The drag is that all vectors for XOPs are in ROM memory, and only three of them (two with some consoles) have usefull values. That's because the GPL interpreter code begins right there. However, there are some eight empty words at the end of the console ROM, TI could just have shifted up the whole stuff and provide us with 4 more XOPs. Oh well, we can do with the first three. Not to mention that some of the following vectors happen to contain usefull values.

Address	XOP	WR	PC	Comments
>0040	0	>280A	>0C1C	Enters the extended GPL card
>0044	1	>FFD8	>FFF8	Very usefull for us
>0048	2	>83A0	>8300	Very usefull for us, but not always present
>004C	3	>1100	>06A0	
>0050	4	>0864	>06A0	
>0054	5	>0864	>C90D	
>0058	6	>8300	>C342	Could be used, although not meant to be so
>005C	7	>D11D	>C180	Dangerous (odd WR)
>0060	8	>DB46	>0402	Pops inside the keyscan routine with wrong WS
>0064	9	>0B60	>83ED	
>0068	10	>0402	>5802	

>006C	11	>011B	>837C	
>0070	12	>0300	>0002	
>0074	13	>0300	>0000	
>0078	14	>D25D	>1105	Pops inside XML >0E with wrong WS
>007c	15	>D109	>09C4	Pops into the ISR (sprite motion) with wrong WS

Now, any WR value that maps to the console ROM (below >2000) is useless as it will result in losing the return address. I suspect that odd workspace addresses may also cause havoc. The same is true for PC values: we don't want to branch to ROM routines. Odd PC values have less importance since the TMS9900 will ignore the least significant bit anyway.

**XOP 0** is hardwired to switch on a peripheral card whose CRU base address should be >1B00, then enters its ROM at address >4028, after having changed the workspace to >2800. My guess is that this card was meant to implement extra GPL opcodes, but I don't think it has ever been released. Note that XOP 0 will not check whether the card is here or not before branching. If there is no such card, the TI-99/4A will crash. Now this is the ideal instruction to use to implement a debugger board...

**XOP 1** is extremely useful for us. All we need to do is to place a B @MYPROG at location >FFF8 and XOP 1 will enter our program. Note that this will preserve the LOAD interrupt vectors at locations >FFFC-FFFF.

**XOP 2** Be careful about XOP 2: some consoles do not support it... If you decide to use it, you should probably place a B @WHERE instruction right at >8300 or soon after, since scratch-pad RAM is a precious resource. Also, having our workspace at >83A0 may disturb the data stack of the GPL interpreter...

**XOP 6** The vectors are not guaranteed to have these values on each and every console...

**XOP 8, XOP 14, XOP 15** land in the middle of various routines in the console ROMs. These routines expect a workspace of >83E0 which won't be the case. What happens then depends on the contents of the workspace in use, of the >83E0 workspace and possibly of other bytes in the scratch-pad. It is not unconceivable that you can come up with a valid combination that would do something useful, but why bother?

Another (admittedly not so useful) trick with the XOP instructions is to use them to build a routine that can be called either with BLWP or with BL. This relies on the fact that XOP opcodes have a value of >2Cxx, which can serve as a valid workspace. See, like this:

```

MYSUB  XOP  R0,1      This is equivalent to DATA >2C40 (WS vector)
      DATA MYSB1    PC vector

MYSB1  ...           Do something
      RTWP           Return to caller (or to BLS)

*-----
* This routine transforms BL calls into BLWP calls
*-----
BLS    MOV  R13,@>2C5A  Put user's workspace pointer in future R13
      MOV  @22(R13),@>2C5C Put return address in future R14
      MOV  R15,@>2C5E  Put user's status in future R15
      MOV  *R14+,@>2C56 Get address from PC vector, in future R11
      LWPI >2C40      Change workspace
      B    *R11       Branch to routine

*-----
* Set up XOP 1
*-----
      AORG >FFF8
      B    @BLS       Entered by XOP 1

      END

```

If the procedure is called with `BLWP @MYSUB` it will treat the word at `MYSUB` as a workspace pointer (of value `>2C40`) and begin execution of the subroutine at `MYSB1` with a workspace of `>2C40`. Return to the caller is performed by a plain vanilla `RTWP`

If it is called with `BL @MYSUB` it executes `XOP 1` which immediately branches to `BLS` with workspace `>FFD8`. `BLS` is a routine that converts `BLs` into `BLWPs`: it gets the PC vector from the data word following the `XOP 1` instruction (in this case `MYSB1`) and puts it in `R11` of workspace `>2C40` (that receives the source argument `R0` upon `XOP` execution). Then it copies the user's workspace pointer and status (saved by `XOP 1`) into `R13` and `R15` of workspace `>2C40`. It gets the return address from the `R11` in the user's workspace and places it in `R14` of workspace `>2C40`. Finally it branches to the called subroutine, which will never be aware of all the above: it can just assumed it was called with `BLWP`, access parameters accordingly and return with `RTWP`.

That's a helluva slow way to call a subroutine, but it might be usefull in cases...

## X instruction

This instruction can be used to simulate another instruction: just place the corresponding code in the source argument of the `X` (`eXecute`) instruction.

Example:

```
LI    R0,>37C3    >37C3 means STCR R3,15
X     R0
```

If the executed instruction has operands, they will be fetched from the words following `X`, which somewhat limitates the usefullness of `X`. If it were to fetch the operands from the words following the operand of the `X` instruction, it would be a wonderfull way to write a debugger: place a memory pointer in `R1` and then do: `X *R1+` Ok,ok it's not that simple: we must trap the jumps and branches, and account for instructions that use `R1`, but you get the idea. Unfortunately, that's not the way `X` works...

Nevertheless, `X` may be usefull to replace a test in a frequently executed loop. For instance:

```
      MOV @CLEAR,R2
      MOV R5,R5      Performs some test
      JEQ SK1        Decide whether to clear or set to one
      MOV @SET,R2     Set to one

SK1   LI    R0,>2000   This loop is executed >2000 times
      LI    R1,BUFFER
LP1   X     R2         Equals CLR *R1+ if R5 was null, SET0 *R1+ otherwise
      DEC   R0
      JNE   LP1
      ...

CLEAR CLR *R1+        Simple alternative to calculating the values
SET   SET0 *R1+       of these instructions to put them in R2
```

We could have written:

```
      LI    R0,>2000
      LI    R1,BUFFER
LP1   MOV   R5,R5      Perform the test inside the loop
      JEQ   SK1
      SET0  *R1+        Set
      JMP   SK2
SK1   CLR   *R1+        Clear
SK2   DEC   R0
      JNE   LP1
      ...
```

But the first way is much faster, since we don't have to repeat the test at each execution of the loop.

## C instruction

Appart for comparison, this instruction can also be used to increment a register by four:

```
C    *Rx+, *Rx+
```

This uses only one word of memory as opposed to the equivalent :


```
INCT Rx
INCT Rx
```

Note that the corresponding CB instruction would increment the register by two, but there is no advantage over a plain vanilla INCT in this case.

## Timing diagrams

### Clock signals

The TMS9900 is meant to be fed a 3 MHz clock signal (that's right, 3 not 30) by the TMS9904 clock generator. This signal comes on 4 different lines, each one being shifted by a quarter of a phase with respect to the previous one. A graphical representation of these signals looks like this:



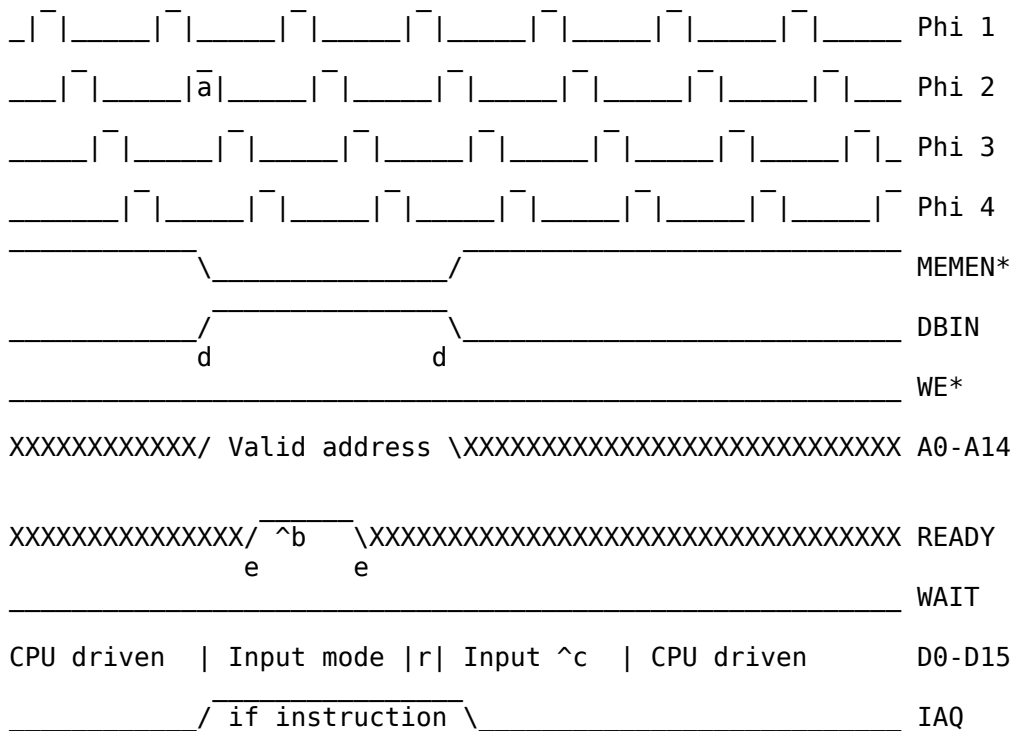
The period of the clock, i.e the time between two pulses in a given phase is 333 nanoseconds for a 3 Megahertz clock (a nanosecond is a billionth of a second). Each pulse is high for about 45nanoseconds, with a rising time of 12 ns and a falling time of 12 ns (note that my graph is not drawn to scale).

Pulses in the next phase are 83 ns behind those in the first phase and there is a 5 ns lag time between the end of one pulse and the start of the corresponding pulse on the next phase. At least, that's what the data manual says, but if you add up durations: 12+45+12+5 you get 74 ns, not 83 ns! So were are the missing 9 nanoseconds? Your guess is mine...


Note that it is possible to crank up clock speed, upto 4 Mhz at least, without risking to fry the TMS9900. Such modifications the the TI-99/4A have been described (including in those pages), and are known to work. Of course any process that relies on execution speed to time an external device (such as disk access) will be messed up...

### Internal memory bus timing

Below are some timing diagrams for the memory bus of the TMS9900. Note that these will be different in the PE-Box, due to the multiplexing of the data bus.

**Read cycle (no wait state)**Notes

- a) The cycle begins and ends on the rising edge of Phi 2 pulses.
- b) Inputs should be ready at least 30 ns before the rising edge of the next Phi 1 pulse.
- c) Inputs should remain valid for at least 10 ns after the falling edge of the Phi 1 pulse.
- d) Propagation delays are at most 30 ns for MEMEN\*, DBIN, WE\* and WAIT.
- e) Propagation delays for all other outputs are at most 40 ns.
- r) Read data

**Write cycle (1 wait state)**Notes


- a) The cycle begins with the rising edge of Phi2.

b) The READY line is tested on the rising edge of the next Phi 1. It should be high at least 40 ns before that time..

c) If it's low, the TMS9900 enters a wait state and

d) retest the READY line at each Phi 1 pulse, until it is high again.

### CRU output (2 bits)



#### Notes

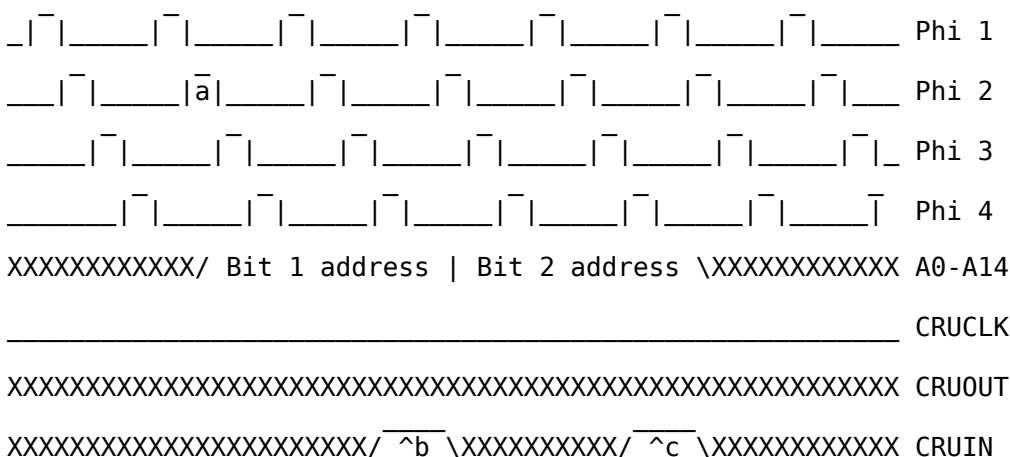
a) The cycle begins on the rising edge of a Phi 2 pulse.

b) A CRUCLK pulse is issued at the next Phi 2 pulse, until the end of the Phi 3 pulse.

c) A similar CRUCLK pulses is issued for each following bit.

d) Propagation delays are at most 30 ns for CRUCLK (40 ns for other outputs).

### CRU input (2 bits)



#### Notes

a) The cycle begins on the rising edge of a Phi 2 pulse.

b) The CRUIN line is sampled at the rising edge of the second Phi 1 pulse following the Phi 2 pulse.

c) Following bits are sampled on the rising edge of every second Phi 1 pulse.

d) No CRUCLK is generated during CRU input.

## Electrical characteristics



## Recommended operating conditions

Parameter	Min	Nom	Max	Unit
Vbb	5.25	-5	-4.75	Volts
Vcc	4.75	5	5.25	Volts
Vdd	11.4	12	12.6	Volts
Vss	-	0	-	Volts
High level input	2.2	2.4	Vcc+1	Volts
Ditto for clock	Vdd-2	-	Vdd	Volts
Low level input	-1.0	0.4	0.8	Volts
Ditto for clocks	-0.3	0.3	0.6	Volts
Free-air temperature	0	25	70	°C

## Electrical characteristics under recommended conditions

Parameter	Test conditions	Min	Nom	Max	Unit
Data bus input current	Vss to Vcc	-	50	100	uAmp
Clock input current	-0.3V to 12.6V	-	25	75	uAmp
Other pins input current	Vss to Vcc	-	1	10	uAmp
High level output voltage	-0.4 mAmp	2.4	-	Vcc	Volts
Low level output voltage	3.2 mAmp	-	-	0.65	Volts
	2.0 mAmp	-	-	0.50	Volts
Supply current from Vbb	-	-	0.1	1	mAmp
Supply current from Vcc	-	-	50	75	mAmp
Supply current from Vdd	-	-	25	45	mAmp
Data bus capacitance	Vbb=-5 f=1Mhz	-	15	25	pF
Clock 1 capacitance	Vbb=5 f=1Mhz	-	100	150	pF
Clock 2 capacitance	Vbb=5 f=1Mhz	-	150	200	pF
Clock 3 capacitance	Vbb=5 f=1Mhz	-	100	150	pF
Clock 4 capacitance	Vbb=5 f=1Mhz	-	100	150	pF
Other input capacitance	Vbb=5 f=1 MHz	-	10	15	pF

Revision 1. 3/25/99 Preliminary, but ok to release

Revision 2. 5/31/99 Tested & debugged examples

Revision 3. 1/4/02 Modified speed calculation examples

Revision 4. 1/15/06 Added example of external instruction use (Cortex computer)

[Back to the TI-99/4A Tech Pages](#)